

# ECE-GY 6143: Introduction to Machine Learning

## Final Exam Solutions, Spring 2019

Prof. Sundeep Rangan

1. *SVM*. Consider an SVM classifier,

$$\hat{y} = \begin{cases} 1 & \text{if } z > 0, \\ -1 & \text{if } z \leq 0, \end{cases} \quad z = \sum_{i=1}^N \alpha_i y_i K(x, x_i) + b,$$

for the kernel function,

$$K(\mathbf{x}, \mathbf{x}_i) := \begin{cases} 1 & \text{if } \|\mathbf{x} - \mathbf{x}_i\|^2 \leq r^2, \\ 0 & \text{if } \|\mathbf{x} - \mathbf{x}_i\|^2 > r^2, \end{cases}$$

where  $r > 0$  is some parameter and  $\|\mathbf{w}\|^2$  denotes the squared norm  $\|\mathbf{w}\|^2 = w_1^2 + w_2^2$ . You are given four training samples.

$i$	1	2	3	4
$x_{i1}$	0	1	0	2
$x_{i2}$	0	0	2	2
$y_i$	1	-1	-1	1

- (a) Find parameters  $r$ ,  $\alpha_i$  and  $b$  such that the SVM classifier makes no errors on the training data. Use at most two non-zero values of  $\alpha_i$ .
- (b) Given the choice of parameters in (a), draw the region of  $(x_1, x_2)$  where the classifier predicts  $\hat{y} = 1$ .
- (c) Complete the following python function `predict` that computes a vector of outputs `yhat` for a data matrix `x`. You must provide the other arguments needed. For full credit, avoid for loops.

```
def predict(X, ...):  
    ...  
    return yhat
```

**Solution:**

(15 points total)




Figure 1: Scatter plot of the data points where the red circles are  $y_i = -1$  and blue are  $y_i = 1$ . The hashed regions are the areas where the classifier will set  $\hat{y} = 1$ .

- (a) (6 points) Take  $\alpha = [1, 0, 0, 1]$ ,  $b = -0.5$  and  $r = 0.5$ . Then,

$$\begin{aligned} z &= \sum_{i=1}^N \alpha_i y_i K(x, x_i) + b \\ &= K(x, x_1) + K(x, x_4) - 0.5. \end{aligned}$$

Since  $r = 0.5$ , we have  $K(x_i, x_j) = 0$  for all training points  $x_i \neq x_j$  since the training points are separated by more than  $r$ . Therefore, at  $x = x_1$  we have

$$z = K(x_1, x_1) + K(x_1, x_4) - 0.5 = 1 + 0 - 0.5 > 0.$$

Similarly, at  $x = x_4$ ,  $z > 0$ . But, at  $x = x_2$ ,

$$z = K(x_2, x_1) + K(x_2, x_4) - 0.5 = 0 + 0 - 0.5 < 0.$$

Similarly, at  $x = x_3$ ,  $z < 0$ . Therefore, the classifier classifies all the points correctly.

- (b) (4 points) See Fig. 1.

- (c) (5 points, 3 points if for loops were used instead of Python broadcasting) One solution is as follows:

```
def predict(X, Xtr, ytr, b, alpha, r):
    # Compute the distances
    # D[i, j] = ||X[i, :] - Xtr[j, :]||^2
    D = np.sum((X[:, None, :] - Xtr[None, :, :])**2, axis=2)

    # Compute the kernel
    K = (D < r**2)

    # Compute the score
```




Figure 2: Green lines: Boundaries of the three hidden units. The hidden units  $u_j^H = 1$  in the interior of the triangle. Blue circles: Points classified to  $\hat{y} = 0$ ; Red circle: Point classified to  $\hat{y} = 1$ .

```
z = K.dot(ytr*alpha) + b
yhat = 2*(z > 0)-1 # Note: We convert to +/- one
return yhat
```

2. *Neural Networks.* Consider a neural network, with  $N_i = 2$  input units,  $\mathbf{x} = (x_1, x_2)$ ,  $N_h = 3$  hidden units and one output unit for binary classification:

$$z_j^H = \sum_{k=1}^{N_i} W_{jk}^H x_k + b_j^H, \quad u_j^H = \begin{cases} 1 & \text{if } z_j^H > 0 \\ 0 & \text{if } z_j^H \leq 0, \end{cases} \quad j = 1, \dots, N_h$$

$$z^O = \sum_{k=1}^{N_h} W_k^O u_k^H + b^O, \quad \hat{y} = \begin{cases} 1 & \text{if } z^O > 0 \\ 0 & \text{if } z^O \leq 0. \end{cases}$$

- (a) Suppose that

$$\mathbf{W}^H = \begin{bmatrix} 0 & -1 \\ 1 & 1 \\ -1 & 1 \end{bmatrix}, \quad \mathbf{b}^H = \begin{bmatrix} 2 \\ -1 \\ 1 \end{bmatrix}$$

For each hidden unit  $j = 1, 2, 3$ , draw the regions of inputs  $(x_1, x_2)$  where  $u_j^H = 1$ .

- (b) Find a vector of output weights  $\mathbf{W}^O$  and bias  $b^O$  such that:

- $\mathbf{x} = (0, 0), (2, 0), (1, 3)$  are classified as  $\hat{y} = 0$ ; and
- $\mathbf{x} = (1, 0.5)$  is classified as  $\hat{y} = 1$ .

**Solution:**

(14 points)

(a) (7 points) We have

$$\mathbf{z}^H = \mathbf{W}^H \mathbf{x} + \mathbf{b}^H = \begin{bmatrix} -x_2 + 2 \\ x_1 + x_2 - 1 \\ -x_1 + x_2 + 1 \end{bmatrix}.$$

Hence,

$$\begin{aligned} u_1^H &= 1 \iff x_2 \leq 2 \\ u_2^H &= 1 \iff x_2 \geq 1 - x_1 \\ u_3^H &= 1 \iff x_2 \geq -1 + x_1 \end{aligned}$$

The boundaries of the three regions are shown in Fig. 2.

(b) (7 points) The points to be classified as  $\hat{y} = 0$  and  $\hat{y} = 1$  are shown, respectively, in blue and red in Fig. 2. Now, we select  $\mathbf{W}^O = [1, 1, 1]$  and  $b^O = -2.5$ . Then,  $z^O > 0$  only when  $u_j^H = 1$  for all  $j$ . This is the triangular region in the Fig. 2, which contains the red point, but all the blue points are outside this region.

3. *Backpropagation.* Consider the model with  $D$ -dimensional inputs  $\mathbf{x} = (x_1, \dots, x_D)$  and  $M$ -dimensional outputs  $\mathbf{y} = (y_1, \dots, y_M)$ , given by

$$\hat{y}_m = \sum_{\ell=1}^L B_{\ell m} z_{\ell}, \quad z_{\ell} = \sum_{j=1}^D x_j A_{j\ell}, \quad m = 1, \dots, M,$$

with parameters  $A_{j\ell}$  and  $B_{\ell m}$ .

- (a) You are given training samples  $\mathbf{x}_i = (x_{i1}, \dots, x_{iD})$  and  $\mathbf{y}_i = (y_{i1}, \dots, y_{iM})$ . Write  $\hat{y}_{im}$  in terms of the inputs  $x_{ij}$ .
- (b) Draw the computation graph between the data, parameters and loss function using the loss function,

$$J = \sum_{i=1}^N \sum_{m=1}^M (y_{im} - \hat{y}_{im})^2.$$

(c) Show how to compute  $\partial J / \partial A_{j\ell}$  from  $\partial J / \partial z_{i\ell}$ .

**Solution:**

(15 points)

(a) (5 points) We have,

$$\hat{y}_{im} = \sum_{\ell=1}^L B_{\ell m} z_{i\ell}, \quad z_{i\ell} = \sum_{j=1}^D x_{ij} A_{j\ell}.$$

(b) (5 points) See Fig. 3.




Figure 3: Computation graph for Problem 3. Parameters are shown in light blue and data in light green.

(c) (5 points) We have

$$\partial z_{il} / \partial A_{jl} = x_{ij}.$$

By chain rule,

$$\frac{\partial J}{\partial A_{jl}} = \sum_i \frac{\partial J}{\partial z_{il}} \frac{\partial z_{il}}{\partial A_{jl}} = \sum_i \frac{\partial J}{\partial z_{il}} x_{ij}.$$

4. *CNN kernels.* A systems has five cameras that each take  $256 \times 192$  dimension gray scale images. A mini-batch of training samples consists of 100 samples, with each sample having the five camera images.

- What is the shape of a tensor  $X$  representing the mini-batch of training data.
- A first layer of a CNN performs a convolution with 20 output channels and  $3 \times 3$  kernels  $W$ . Write the equation relating the input  $X$  and output  $Z$ . What is the shape of  $W$  and output  $Z$ . Assume the convolution is performed on the valid pixels.
- Describe a possible kernel  $W$  that detects a difference between camera 2 and camera 3, but ignores cameras 0,1, and 4. The difference output should be in output channel 0.

**Solution:**

(14 points)

(a) (4 points) Take the shape (sample,row,height,camera) for (100,256,192,5).

(b) (5 points) We have

$$Z[i, j_1, j_2, m] = \sum_{k_1, k_2} X[i, j_1 + k_1, j_2 + k_2, n] W[k_1, k_2, n, m].$$

The kernel  $W$  will have shape (3, 3, 5, 20). Since the output is computed on the valid pixels,  $Z$ , and the kernels are  $3 \times 3$ , will have shape  $(100, 256 - 3 + 1, 192 - 3 + 1, 20) = (100, 254, 190, 20)$ .

(c) (5 points) Take

$$W[:, :, n, 0] = \begin{cases} \frac{1}{9} & \text{if } n = 2 \\ -\frac{1}{9} & \text{if } n = 3 \\ 0 & \text{else.} \end{cases}$$

In this way, the kernels take the average of cameras 2 and 3 and subtract the two averages.

5. *CNN sub-sampling.* Consider a 1D convolution following by sub-sampling,

$$z[j] = \sum_k w[k]x[j+k], \quad u[m] = z[sm],$$

for some stride parameter  $s > 0$ . That is,  $u$  takes every  $s$  samples of  $z$ .

- (a) If  $x$  is 250 milliseconds of audio sampled at 20 kHz, how many samples are in  $x$ ?
- (b) If  $x$  has length 1000,  $w$  has length 10 and  $s = 4$ , how many output samples are in  $u$  assuming the convolution is only computed on the valid samples.
- (c) Given a gradient  $\partial J / \partial u[m]$ , how do you compute  $\partial J / \partial w[k]$ ?

**Solution:**

(14 points)

- (a) (4 points) The number of samples are  $(20)(250) = 5000$ .
- (b) (4 points) Before sub-sampling, there are  $1000 - 10 + 1 = 991$  samples. So, there will be  $\lfloor 991/4 \rfloor = 247$  samples
- (c) (6 points) Given a gradient  $\partial J / \partial u[m]$ , how do you compute  $\partial J / \partial w[k]$ ?

We have,

$$u[m] = z[sm] = \sum_k w[k]x[sm+k],$$

so,

$$\frac{\partial u[m]}{\partial w[k]} = x[sm+k].$$

By chain rule,

$$\frac{\partial J}{\partial w[k]} = \frac{\partial J}{\partial u[m]} \frac{\partial u[m]}{\partial w[k]} = \frac{\partial J}{\partial u[m]} x[sm+k].$$

6. *PCA.* You are given python arrays with PCA data:

- `mu`: A 100-dim vector representing the mean of the data; and
  - `v`: A (100,5) array with the 5 top PCs of the data.
  - `lam`: A 5-dim vector of eigenvalues.
- (a) Given a vector `z` of PC coefficients, write a few lines of python code to reconstruct the data `x`.

- (b) Now, suppose you are given only the first 80 of the 100 coefficients of a vector  $\mathbf{x}$ . Write a few lines of python code to estimate the PCA coefficients  $\mathbf{z}$ . You can assume you have a function,

```
w = lstsq(A,b) # Solves the least-squares solution to b=A.dot(w)
```

- (c) Continuing in part (b), how would you estimate the remaining 20 coefficients of  $\mathbf{x}$ ?

**Solution:**

(14 points)

- (a) (4 points) One solution is:

```
xhat = mu + V.dot(z)
```

- (b) (6 points) We know that  $x \approx \mu + Vz$ , so the first 80 components are,  $x[:80] \approx \mu[:80] + V[:80,:]\mathbf{z}$ . We can then solve via least squares:

```
z = lstsq(V[:80,:], x[:80]-mu[:80])
```

- (c) (4 points) We can recover the last 20 components:

```
xhat[80:] = V[80:,:].dot(z) + mu[80:]
```

7. *K-means*. You are given five data samples:

$i$	1	2	3	4	5
$x_{i1}$	0	1	0	2	2
$x_{i2}$	0	0	1	2	3

- (a) Draw the five points.
- (b) Starting with  $K = 2$  cluster centers at  $(0,0)$  and  $(1,0)$ , what are the cluster assignments and new cluster centers after one iteration of *K-means*?
- (c) Suppose you want to use clustering for outlier detection. You find cluster means  $\mu_i$ ,  $i = 1, \dots, K$  on the training data. Then, given a new data  $\mathbf{x}$  and a threshold  $t$ , you declare  $\mathbf{x}$  an outlier if  $\|\mathbf{x} - \mu_i\| \geq t$  for all  $i$ . Complete the following function to implement the outlier detection on a matrix of data  $\mathbf{x}$ . The output is `out[i]=1` if the sample  $\mathbf{x}[i,:]$  is an outlier, and `out[i]=0` otherwise. You must specify the other inputs of your function. Avoid for loops for full credit.

```
def outlier_detect(X, ...):
    ...
    return out
```

**Solution:**

- (a) (4 points) The five points are shown in Fig. 4.
- (b) (4 points) The starting clusters are  $\mu_1 = (0,0)$  and  $\mu_2 = (1,0)$ . The points closest




Figure 4: Data points for  $K$ -means clustering 7.

to these cluster centers are:

$$C_1 = \{(0, 0), (0, 1)\},$$

$$C_2 = \{(1, 0), (2, 2), (2, 3)\}.$$

The new cluster centers are

$$\mu_1 = \frac{1}{2} [(0, 0) + (0, 1)] = (0, 0.5),$$

$$\mu_2 = \frac{1}{3} [(1, 0) + (2, 2) + (2, 3)] = \left(\frac{5}{3}, \frac{5}{3}\right).$$

(c) (4 points) One solution to this (in python) is as follows:

```
def outlier_detect(x,mu,t):
    # Compute squared distance to the clusters
    # dsq[i,j] = distance from x[i,:] to mu[j,:]
    dsq = np.sum((x[:,None,:] - mu[None,:,:])**2, axis=2)

    # Find the minimum distance
    dmin = np.min(dsq, axis=1)

    # Declare an outlier if minimum distance
    # exceeds threshold
    out = (dmin > t**2)
    return out
```